Relaxed Softmax: Efficient Confidence Auto-Calibration for Safe Pedestrian Detection

Lukas Neumann **Andrew Zisserman** **Andrea Vedaldi**
Department of Engineering Science
University of Oxford
{lukas,vedaldi,az}@robots.ox.ac.uk

Abstract

As machine learning moves from the lab into the real world, reliability is often of paramount importance. The clearest example are safety-critical applications such as pedestrian detection in autonomous driving. Since algorithms can never be expected to be perfect in all cases, managing reliability becomes crucial. To this end, in this paper we investigate the problem of learning in an end-to-end manner object detectors that are accurate while providing an unbiased estimate of the reliablity of their own predictions. We do so by proposing a modification of the standard softmax layer where a probabilistic confidence score is explicitly pre-multiplied into the incoming activations to modulate confidence. We adopt a rigorous assessment protocol based on reliability diagrams to evaluate the quality of the resulting calibration and show excellent results in pedestrian detection on two challenging public benchmarks.

1 Introduction

Deep neural networks have been very successful at addressing many real-world classification, detection and recognition problems. As a result, they have already moved from being a subject of research to being deployed "in the wild", where their decisions impact everyday life.

Among such applications, some, such as pedestrian detection in autonomous driving, are critical for safety. In such cases, the ability of networks to make reliable decisions is of paramount importance. As new architectures and components are introduced, research seems to be focused on improving accuracy and speed of networks *on average*. Yet, far less attention has been given to determining *when* predictions can be trusted and when they cannot. In other words, neural networks are good at providing answers which are correct most (but not all) of the times, but not good at telling when such answers can be trusted, or when, instead, they amount to little more than an educated guess.

Having a measure of the network's prediction reliability is essential in many applications. While the strength of activations in deep neural networks is often correlated with the network's confidence, such values can hardly be mapped in a systematic manner to a meaningful and verifiable measure of confidence. Consider an application such as pedestrian detection using a camera mounted on the front of a car. In most situations, the visual information is clear and the network can determine with confidence whether a pedestrian is contained in the field of view of the camera or not. However, in some cases viewing conditions are much poorer. A good example is pedestrian detection at night (Figure 1), in which poor, unusual, and rapidly changing illumination, shadows, reduced color sensitivity, and tendency of colors to fuse in the background make reliable detection exponentially harder.

In such cases, we would like to know when the network is not able to make an accurate determination of the possible presence of pedestrians in front of the car, for example in order to take precautionary action such as slowing down, or relying more on other sensors. Instead, as we confirm empirically in this paper, neural networks do not hedge their bets, and vote with exceeding confidence for one case (pedestrian present) or the other (not present) even when the evidence in the data is genuinely scarce or ambiguous. Our first contribution is, in fact, to assess state-of-the-art pedestrian detectors in terms of the quality of their confidence estimates using reliability diagrams and the calibration error (CE) metrics.

Our second contribution is of a technical nature and amounts to a simple but effective modification of the standard softmax layer used in neural networks that significantly improves its self-calibration properties. This layer, which we call **relaxed softmax**, has nearly the same complexity as softmax, requiring the network to

2018 NIPS Workshop on Machine Learning for Intelligent Transportation Systems.

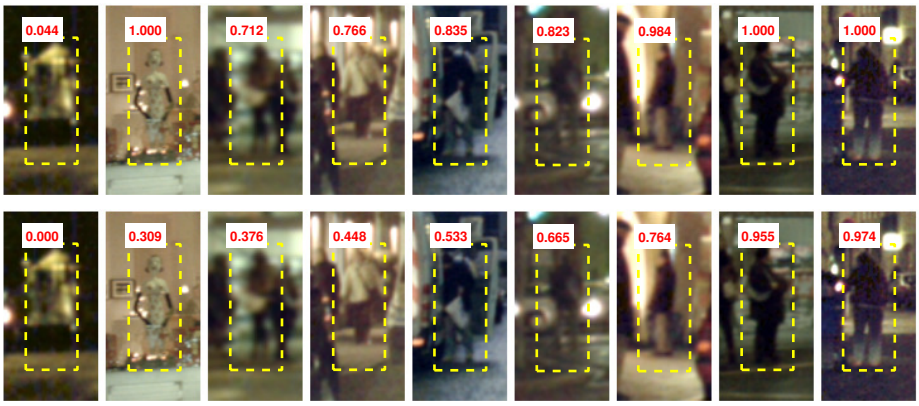

Figure 1: **Decision confidence in safety-critical applications (NightOwls data)**. Top: detection probability estimated by a state-of-the-art neural network detector applied to pedestrians. The probabilities (number in red) do not reflect well the amount of evidence, with overconfident detections in ambiguous cases (*e.g.*, the statue in 2nd column above). Bottom: *relaxed softmax* (bottom row) provides *calibrated* output that significantly better reflects the detection ambiguity.

estimate just one more scalar output per sample, which is of negligible computational impact in all practical architectures.

We train relaxed softmax using the standard cross-entropy loss that is already used with vanilla softmax, which should result in both being properly calibrated. However, as found by other authors [24, 8] and as we confirm in the experiments, reliable calibration does not occur for softmax. There are two possible opposite explanations: overfitting, as the model becomes overly confident on the training data, and underfitting, as the architecture may be unable to easily represent confidence.

We show empirically that both effects are present. Overfitting is present because we can improve calibration by simply rescaling all scores by a constant multiplicative factor tuned on a validation set. Underfitting is present as well, as relaxed softmax, which is used as a drop-in replacement for softmax, does significantly improve the calibration profile of the emitted probabilities and achieves a significant reduction in CE. While we show this in the context of object detection, and safety-critical pedestrian detection in particular, we expect similar benefit to apply anywhere softmax is used.

The rest of the paper is organized as follows. Section 2 summarizes the work most related to ours, Section 3 discusses our method in detail, Section 4 assess it empirically in challenging pedestrian detection scenarios, and Section 5 summarizes our findings.

2 Related Work

Modeling Uncertainty. Compared to the vast amount of literature focused on improving the accuracy [25, 10, 27, 12] and speed [22, 16, 21] of deep neural networks, not much research attention has been focused on studying the confidence of their outputs. Several authors studied the reliability of classical logistic regression [1] and SVM [13, 29, 7] classifiers. In the context of deep networks, Gal and Ghahramani [8] study Dropout [23] and its relation to Gaussian processes. Kendall and Cipolla [14] estimate uncertainty in a camera pose regression problem, exploiting Monte Carlo sampling. Novotny et al. [20] use network uncertainty as a mean to learn 3D geometry predictions from videos in an unsupervised manner. More recently, Kendall and Gal [15] model uncertainty by making logits in a softmax follow a Gaussian (Laplacian) distribution and train the network by optimizing a stochastic loss through Monte Carlo integration. None of the aforementioned methods however try to explicitly output model confidence, nor they systematically evaluate the network's calibration error.

Recently, Guo et al. [9] studied the problem of deep networks calibration, but only focused on parametric and non-parametric calibration of pre-trained models that have no notion of uncertainty in their own right. Our model, in contrast, is specifically trained to infer its own uncertainty. Additionally, they only focus on simple classification problems based on the expected calibration error metric, which is not suitable for detection problems (see Section 3.3).

Pedestrian Detection. Pedestrian detection is one of the most intensely studied subfield in object detection, because of its importance in real-world applications. All current state-of-thea-art methods are derived from the Faster R-CNN detector of Ren et al. [22], but they differ in the way they address the issue of detecting small pedestrians. **?**] and Cai et al. [3] use a specialized multi-scale networks, whereas other methods [28, 31, 2] exploit a separate classification network which operates on higher-resolution features or directly on the source image.

The current state-of-the-art pedestrian detection method of Brazil et al. [2] also falls within this category, as their method is a two-stage pedestrian detector combined with an attention-based mechanism coming from a coarse semantic segmentation branch. In the first stage, region proposals are generated by a Region Proposal Network (RPN), and in the second stage, the proposals are classified as pedestrian or a background using a separate network that operates directly on image crops.

3 Method

This section describes our approach for end-to-end calibration of object detectors. We start by reviewing the general architecture of modern detectors based on convolutional neural networks (Section 3.1) and then introduce a relaxed softmax layer that can significantly improve the quality of the detection confidence estimated by the neural network (Section 3.2).

3.1 Background: Object Detectors

Most modern detectors, including variants specialized for the detection of pedestrians [28, 30, 31, 2], are based on the *propose & verify* paradigm. Namely, given an image I, a mechanism such as a specialized neural network [22], an ad-hoc algorithm [26], or even simply the use of a fixed pool [16], is used to generate a shortlist of candidate image regions R_i, $i = 1, \dots, N$ that are likely to contain occurrences of the objects of interest. Then a neural network $\phi(R_i; I)$ is tasked with mapping each region to a vector $(z_{i,1}, \dots, z_{i,K+1}) \in \mathbb{R}^K$ of scores, one for each possible class identity $1, \dots, K$. Scores are converted into probabilities by using the softmax operator σ, resulting in a normalized belief value $\sigma_i(j)$ for each hypothesis $j = 1, \dots, K$:

$$\sigma_i(j; z_i) = \frac{\exp(z_{i,j})}{\sum_{k=1}^{K} \exp(z_{i,k})}. \tag{1}$$

In this manner, the detection problem is reduced to a standard classification problem. As commonly done in deep networks, after converting the scores produced by Φ into probabilities by using the softmax operator, scores are fitted to the ground truth label $y_i \in \{1, \dots, K+1\}$ of each region by minimizing the cross-entropy loss:

$$\mathcal{L}_{\text{cls}}(z_i, y_i) = -\sum_i \log \sigma_i(y_i; z_i). \tag{2}$$

The ground-truth label y_i of each candidate region R_i is determined to be positive or negative based on the Intersection-over-Union (IoU) of R_i with the ground truth object annotations, which must therefore be available. This process, which effectively transfers labels from the annotated regions to the candidate regions, is required because none of the candidates will in general match exactly the ground-truth annotations.

After training, given a new image, the neural network Φ can associate to each candidate region R_i a class y_i with confidence s_i as the maximum of the score vector $\sigma_i(j)$, $j = 1, \dots, K$:

$$y_i = \operatorname*{argmax}_j \sigma_i(j; z), \qquad s_i = \sigma_i(y_i; z). \tag{3}$$

Each candidate box is then associated with its position, size and a score, which determines how likely the candidate box is to contain the object of interest (pedestrian in our application) or not. During inference, only the boxes with the highest score are kept, using non-maxima suppression to filter redundant detections, and candidate boxes with a score above a certain threshold (or the top $M \ll N$ boxes with the highest score) are output by the detector.

3.2 Output Calibration using Relaxed Softmax

Even though the softmax operator σ by definition ensures that the output score s_i are in the interval $(0, 1)$, it has been empirically shown [24, 8] that it does not result in a good estimator of the posterior detection probability $p(y_i|I)$. In other words, the softmax output does not reflect well the model confidence; in practice, deep networks tend to be overconfident in their predictions (see Figure 2a), *i.e.* the scores associated with the detections tend to be higher than the true probability of a correct detection.

Next, we propose a method that encourages the model output s_i to be as close as possible to the posterior $p(y_i|I)$. Based on the observation that networks tend to output high scores even for hard examples, where in principle the confidence should be lower, we let our model express its uncertainty by introducing an additional parameter to the softmax operator that lets the network "soften" its predictions for particular samples. We build on the model of **temperature scaling** [11, 9] for its simplicity, where the classification output is given as

$$\sigma_i(j; z_i, T_i) = \frac{\exp(\frac{z_{i,j}}{T_i})}{\sum_{k=1}^{K} \exp(\frac{z_{i,k}}{T_i})} \tag{4}$$

For $T_i > 1$, this allows the model to "soften" its predictions for the sample i and with $T_i \to \infty$ the score for all classes is identical. For $T_i = 1$, we obtain the original softmax operator (1).

Next, we modify eq. (4) to obtain our *relaxed softmax* layer. The first consideration is to allow the network to learn to self-adjust the temperature *on a sample-by-sample basis*, therefore replacing a global parameter T with a sample-dependent temperature T_i. Equation (4) however is not well-suited for loss optimization

Eq (4) already us

through SGD due to numerical instabilities when $T_i \to 0$. We therefore propose an equivalent formulation, which we refer to as **relaxed softmax**, which is more numerically stable and can therefore easily be plugged in to standard loss optimization frameworks:

$$\hat{\sigma}_i(j; z_i, \alpha_i) = \frac{\exp(\alpha_i z_{i,j})}{\sum_{k=1}^{K} \exp(\alpha_i z_{i,k})}, \qquad \alpha_i := \frac{1}{T_i}. \tag{5}$$

In this model, the network $\Phi(R_i; I)$ predicts a vector $(z_{i,1}, \ldots, z_{i,K}, \alpha_i)$ for each candidate box R_i, and the candidate box is then associated with a score $\hat{\sigma}_i$ using eq. (5). Note that from computational perspective, this is equivalent to predicting $K + 1$ scalar scores instead of K, which is a negligible cost.

We think, that because in our model the network can in fact improve its training objective by outputting "not so sure" decisions ($T > 1$) for ambiguous samples, it is not forced to be over-confident like the standard models where even the ambiguous samples are required by the loss function to output 1.0 confidence. And on contrary, for non-ambiguous samples, the training objective is still maximized by setting $T = 1$, because this puts the probability mass entirely behind the correct class. These two opposing effects, we believe, help the model to become more calibrated, because it allows it to learn and infer how ambiguous the (training) samples are.

3.3 Evaluation Metrics

In order to visually quantify the quality of the calibration of a probabilistic classifier, we adopt the **Reliability Diagrams** of [4, 19, 9]. A reliability histogram shows the score s_i generated by a method along the abscissa, plotting it against the actual probability of correct detection $p(y_i|I)$. Since the true posterior $p(y_i|I)$ is unknown, we empirically approximate the posterior $\hat{p}(y_i|I)$ as a fraction of correctly classified *test* samples, and compare the network's output s_i to the latter. For example, given 100 samples, each with a score $s_i = 0.6$, we would expect 60 samples to be correctly classified and 40 samples to be classified incorrectly.

In the diagram, if the bin which corresponds to the output score 0.6 contains 40% correctly classified samples then the model is *overconfident*, and in contrast, if it contains 80% the model is *conservative*. The desired outcome therefore is that in each bin the method's output score corresponds to the true ratio of correct samples, which is when refer to the method as **calibrated**.

Formally, given the method's output as a set of pairs (s_i, y_i) of predicted scores s_i and labels y_i, the corresponding set of ground truth labels $\hat{y}_i$, and the number of bins M, the *accuracy* A_m is defined as

$$A_m = \frac{1}{|\mathcal{B}_m|} \sum_{i \in \mathcal{B}_m} \mathcal{I}(y_i = \hat{y}_i), \quad \text{where} \quad \mathcal{B}_m = \left\{ i : \frac{m-1}{M} < s_i \leq \frac{m}{M} \right\} \tag{6}$$

where $\mathcal{I}$ is the indicator function. We also define an *average score* S_m for the bin m as

$$S_m = \frac{1}{|\mathcal{B}_m|} \sum_{i \in \mathcal{B}_m} s_i \tag{7}$$

A reliability diagram has S_m on the abscissa and A_m on the ordinate.

Given the reliability diagram (S_m, A_m), we introduce the **Average Calibration Error** (ACE) measure as the primary metric for object detection calibration, as the existing Expected Calibration Error (ECE) metric is not suitable for object detectors (see below). ACE measures the average absolute difference between the score and the accuracy in all bins as

$$\text{ACE} = \frac{1}{M^+} \sum_{m} |S_m - A_m| \tag{8}$$

where M^+ is the number of non-empty bins. ACE therefore assigns an equal weight to each bin, which is desirable for safety-critical applications where we want to measure the total deviation between the predicted score and the accuracy, irrespective to how frequently the object can appear in the image.

We also measure output calibration using the **Expected Calibration Error** (ECE) [17, 9], which weights the error based on the number of samples in each bin:

$$\text{ECE} = \sum_{m} \frac{|\mathcal{B}_m|}{n} |S_m - A_m| \tag{9}$$

This metric however is not very well-suited for the evaluation of object detectors, because there are typically many low-confident detections with a score close to 0, which gives disproportionate weight (typically more than 95%) to the first bin, so the resulting ECE value is then mostly given by the error in the low-confidence predictions.

Finally, we report the **Maximum Calibration Error** (ECE) [17, 9], as this is also very relevant for robust object detection, because it empirically measures the maximal error a method can make in the probability estimation. ECE is given by:

$$\text{MCE} = \max_{m} |S_m - A_m| \tag{10}$$

Not sure if I get this expl

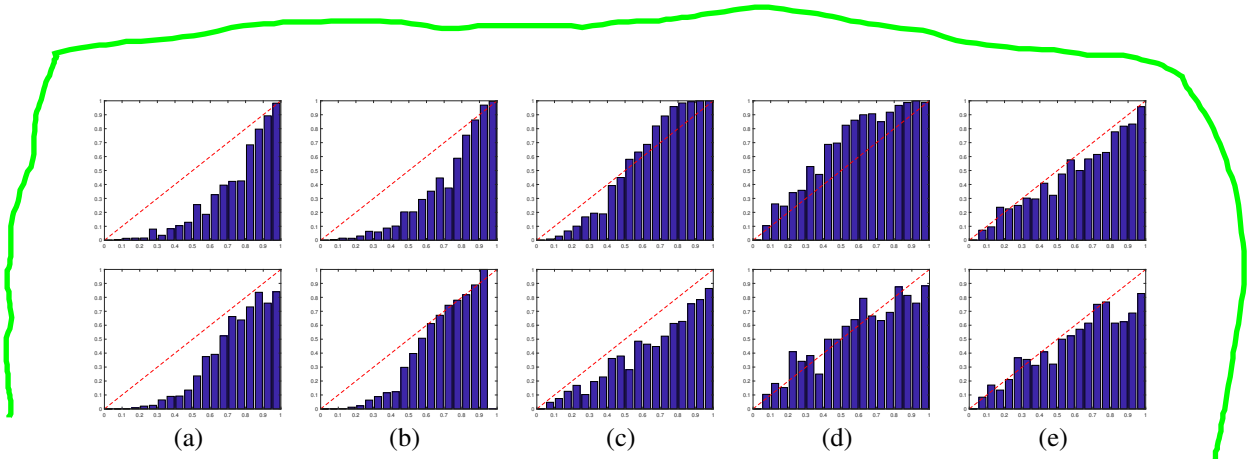

(a) (b) (c) (d) (e)

Figure 2: Reliability Diagrams of different classification output formulations on the Caltech (top) and NightOwls (bottom) dataset. (a) softmax, (b) softmax w/ lin. scaling, (c) softmax with temp. scaling, (d) relaxed softmax and (e) rel. softmax w/ lin. scaling

4 Experiments

In order to compare the traditional softmax to the introduced *relaxed softmax* formulation, we evaluate their properties on two large-scale public benchmarks, using the following experimental setup: We use the state-of-the-art pedestrian detector SDS-RCNN [2] and train two models from scratch — the first model uses softmax (same as [2]), whereas the second model uses *relaxed softmax* for pedestrian/background classification in both stages of the SDS-RCNN. For each model, we also experiment with a subsequent parametric calibration of its output s_i, using the *linear* s_i' and *temperature* scaling s_i'' of Guo et al. [9], given by:

$$s_i' = \min\{1, \beta s_i\} \tag{11}$$

$$s_i'' = \frac{\exp(\frac{z_i}{T})}{\sum_{k=1}^{K} \exp(\frac{z_i}{T})} \tag{12}$$

where β (respectively T) is a single global hyper-parameter, whose optimal value for the given trained model was found on the validation subset by minimizing the Average Calibration Error (ACE). When it is used, linear/temperature scaling is tuned on the validation data after learning the model on the training data.

Next, we show results on the Caltech and NightOwls datasets, comparing the calibration of the output scores s_i, s_i' and s_i'' on the respective *testing set* using the metrics of Section 3.3.

4.1 Caltech Dataset

The Caltech dataset [6, 5] is the most-commonly used dataset for pedestrian detection. For training, we used the same setting as in [2]: the learning rate is set to 10^{-3} and then it is dropped by a factor 10 after every 60,000 iterations. We train the network by minimizing eq. (5) wrt. to α_i and $z_{i,j}$ using vanilla SGD for 180,000 iterations.

We observe that the network which uses softmax is over-confident (see Figure 2a), as the output score is higher that the real ratio of true positives. We also see that linear scaling (11) has little effect on output calibration (see Figure 2b) and that the temperature scaling (12) causes the network to be calibrated only in the center of the $(0, 1)$ confidence interval, but to underestimate number of true positives for scores above 0.5 and inversely to over-estimate the number of true positives below 0.5 (see Figure 2c). This is because temperature scaling only relies on a single global parameter, and therefore the only way to improve the calibration is globally increasing the entropy of the scores distribution [9].

Using the *relaxed softmax* formulation, on the other hand, causes the network to be monotonically more conservative by under-estimating the number true positives in each bin (see Figure 2d) — we think this under-estimation is caused by the fact there are not that many ambiguous training samples in the dataset. Thanks to the monotonicity of the difference, simple linear scaling (11) however makes sure the network becomes almost perfectly calibrated (see Figure 2e). As a result, the method output score can be interpreted as the model confidence.

Quantitatively, the *relaxed softmax* has lower Average Calibration Error (ACE), Expected Calibration Error (ECE) and Maximum Calibration Error (MCE) than softmax and softmax with linear scaling, but has higher ACE than softmax with temperature scaling [9]. *Relaxed softmax* with linear scaling, however, outperforms all other formulations by a significant margin, lowering the ACE 4 times and MCE 3 times when compared to the commonly used softmax (see Table 1).

The average miss rate (using the "reasonable" setting of [6]) of *relaxed softmax* is however slightly worse than standard softmax — we speculate that this is because in training the network "gives up" on certain hard samples, or gives them a lower confidence score, which also negatively impacts the average miss rate, as

Table 1: Average Calibration Error (ACE), Expected Calibration Error (ECE), Maximal Calibration Error (MCE) and average Miss Rate of the SDS-RCNN pedestrian detector on the Caltech dataset, using different classification output formulations.

Classification Output	ACE	ECE	MCE	Miss Rate
Softmax	20.62 %	1.86 %	38.55 %	**10.17** %
Softmax + linear scale	18.52 %	1.81 %	35.26 %	**10.17** %
Softmax + temperature scale [9]	9.58 %	1.95 %	18.49 %	**10.17** %
Relaxed Softmax	14.72 %	0.38 %	29.95 %	13.26 %
Relaxed Softmax + linear scale	**5.62** %	**0.13** %	**14.73** %	13.26 %

Table 2: Average Calibration Error (ACE), Expected Calibration Error (ECE), Maximal Calibration Error (MCE) and average Miss Rate of the SDS-RCNN pedestrian detector on the NightOwls dataset.

Classification Output	ACE	ECE	MCE	Miss Rate
Softmax	17.67 %	4.34 %	33.99 %	19.85 %
Softmax + linear scale	15.37 %	3.75 %	30.08 %	19.85 %
Softmax + temperature scale [9]	12.09 %	0.12 %	24.28 %	19.85 %
Relaxed Softmax	**7.56** %	**0.07** %	**18.84** %	20.62 %
Relaxed Softmax + linear scale	7.74 %	**0.07** %	24.58 %	20.62 %

more false positives will be ranked higher than such samples. Note that the miss rate is not affected by the subsequent parametric calibration, as both transformations are a monotonically increasing function. We can say that the network is slightly worse at guessing, but it is more reliable, which is likely a good trade off in applications.

4.2 NightOwls Dataset

The NightOwls dataset [18] contains 279,000 night images recorded in 3 countries by an industry-standard camera mounted onto a moving car. The dataset is captured at night, at dusk or at dawn, it contains different seasons and weather conditions and is therefore significantly more challenging than the Caltech dataset. It also by definition contains more ambiguous scenarios than the Caltech dataset (see Figure 1), because of low contrast and severe motion blur, caused by high exposure times of the camera. We train again train the SDS-RCNN detector [2] on the training subset for 300,000 iterations using the learning rate of 10^{-3}, which is dropped by the factor of 10 after every 100,000 iterations, and we evaluate the output calibration on the testing subset.

We observe that softmax is again over-confident in its predictions, and unlike for the Caltech dataset, the subsequent calibration on the validation set using linear or temperature scaling does not seem to help (see Figure 2a-c). In contrast, the proposed *relaxed softmax* is significantly better calibrated, even without any subsequent output scaling on the validation set (see Figure 2d and Table 2). We suggest this is because there are more ambiguous training samples in the NightOwls than in the Caltech dataset, which allows the network to learn from samples distributed over the whole $(0, 1)$ confidence interval.

The linear scaling (11) of the output scores then actually makes the calibration error marginally worse (see Figure 2e), which we speculate is due to subtle differences in the data statistics between the validation and the testing set.

5 Conclusion

Having a confidence measure that accurately reflects a neural network's prediction reliability is a crucial requirement for many applications, especially in safety-critical scenarios such as autonomous driving. In this paper, we presented a framework for the systematic assessment of detectors in terms of the quality of their confidence estimates and we showed that current state-of-the-art pedestrian detection methods are over-confident in their predictions when the evidence in the data is genuinely scarce or ambiguous.

Additionally, we proposed a simple but effective modification of the standard softmax layer called *relaxed softmax*, that significantly improves the detectors self-calibration properties. To achieve this, the network has to estimate just one more scalar output per sample, which is of negligible computational impact in all practical architectures. We evaluated its (Average) Calibration Error on two large-scale pedestrian detection datasets, and demonstrated that *relaxed softmax* leads to better output calibration than the pre-existing methods.

We note that relaxed softmax is a drop-in replacement that can be used in any scenario where softmax and cross entropy are used in learning deep neural networks. While we did not yet experiment with other scenarios, the benefits we have observed are likely to transfer to many other cases as well.

How significant is this drop in mi

References

[1] Aayush Bansal, Ali Farhadi, and Devi Parikh. Towards transparent systems: Semantic characterization of failure modes. In *European Conference on Computer Vision*, pages 366–381. Springer, 2014.

[2] Garrick Brazil, Xi Yin, and Xiaoming Liu. Illuminating pedestrians via simultaneous detection & segmentation. In *ICCV*, 2017.

[3] Zhaowei Cai, Quanfu Fan, Rogerio S Feris, and Nuno Vasconcelos. A unified multi-scale deep convolutional neural network for fast object detection. In *European Conference on Computer Vision*, pages 354–370. Springer, 2016.

[4] Morris H DeGroot and Stephen E Fienberg. The comparison and evaluation of forecasters. *The statistician*, pages 12–22, 1983.

[5] P. Dollár, C. Wojek, B. Schiele, and P. Perona. Pedestrian detection: A benchmark. In *CVPR*, June 2009.

[6] Piotr Dollár, Christian Wojek, Bernt Schiele, and Pietro Perona. Pedestrian detection: An evaluation of the state of the art. *PAMI*, 34, 2012.

[7] Abhishek Dutta, Raymond Veldhuis, and Luuk Spreeuwers. Predicting face recognition performance using image quality. *arXiv preprint arXiv:1510.07119*, 2015.

[8] Yarin Gal and Zoubin Ghahramani. Dropout as a bayesian approximation: Representing model uncertainty in deep learning. In *international conference on machine learning*, pages 1050–1059, 2016.

[9] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q Weinberger. On calibration of modern neural networks. In *International Conference on Machine Learning*, pages 1321–1330, 2017.

[10] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778, 2016.

[11] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.

[12] Jie Hu, Li Shen, and Gang Sun. Squeeze-and-excitation networks. *arXiv preprint arXiv:1709.01507*, 2017.

[13] N. Jammalamadaka, A. Zisserman, M. Eichner, V. Ferrari, and C. V. Jawahar. Has my algorithm succeeded? an evaluator for human pose estimators. In *European Conference on Computer Vision*, 2012.

[14] Alex Kendall and Roberto Cipolla. Modelling uncertainty in deep learning for camera relocalization. In *Robotics and Automation (ICRA), 2016 IEEE International Conference on*, pages 4762–4769. IEEE, 2016.

[15] Alex Kendall and Yarin Gal. What uncertainties do we need in bayesian deep learning for computer vision? In *Advances in Neural Information Processing Systems*, pages 5580–5590, 2017.

[16] Wei Liu, Dragomir Anguelov, Dumitru Erhan, Christian Szegedy, Scott Reed, Cheng-Yang Fu, and Alexander C Berg. Ssd: Single shot multibox detector. In *European conference on computer vision*, pages 21–37. Springer, 2016.

[17] Mahdi Pakdaman Naeini, Gregory F Cooper, and Milos Hauskrecht. Obtaining well calibrated probabilities using bayesian binning. In *AAAI*, pages 2901–2907, 2015.

[18] Lukas Neumann, Michelle Karg, Shanshan Zhang, Christian Scharfenberger, Eric Piegert, Sarah Mistr, Olga Prokofyeva, Robert Thiel, Andrea Vedaldi, Andrew Zisserman, and Bernt Schiele. NightOwls: A pedestrians at night dataset. *14th Asian Conference on Computer Vision (ACCV 2018)*, to appear.

[19] Alexandru Niculescu-Mizil and Rich Caruana. Predicting good probabilities with supervised learning. In *Proceedings of the 22nd international conference on Machine learning*, pages 625–632. ACM, 2005.

[20] David Novotny, Diane Larlus, and Andrea Vedaldi. Learning 3d object categories by looking around them. In *The IEEE International Conference on Computer Vision (ICCV)*, October 2017.

[21] Joseph Redmon, Santosh Divvala, Ross Girshick, and Ali Farhadi. You only look once: Unified, real-time object detection. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 779–788, 2016.

[22] Shaoqing Ren, Kaiming He, Ross Girshick, and Jian Sun. Faster R-CNN: Towards real-time object detection with region proposal networks. In *Advances in neural information processing systems*, pages 91–99, 2015.

[23] Nitish Srivastava, Geoffrey E Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. Dropout: a simple way to prevent neural networks from overfitting. *Journal of machine learning research*, 15(1):1929–1958, 2014.

[24] Christian Szegedy, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian Goodfellow, and Rob Fergus. Intriguing properties of neural networks. *arXiv preprint arXiv:1312.6199*, 2013.

[25] Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Sermanet, Scott Reed, Dragomir Anguelov, Dumitru Erhan, Vincent Vanhoucke, Andrew Rabinovich, et al. Going deeper with convolutions. Cvpr, 2015.

[26] Jasper RR Uijlings, Koen EA Van De Sande, Theo Gevers, and Arnold WM Smeulders. Selective search for object recognition. *International journal of computer vision*, 104(2):154–171, 2013.

[27] Saining Xie, Ross Girshick, Piotr Dollár, Zhuowen Tu, and Kaiming He. Aggregated residual transformations for deep neural networks. In *Computer Vision and Pattern Recognition (CVPR), 2017 IEEE Conference on*, pages 5987–5995. IEEE, 2017.

[28] Liliang Zhang, Liang Lin, Xiaodan Liang, and Kaiming He. Is faster r-cnn doing well for pedestrian detection? In *European Conference on Computer Vision*, pages 443–457. Springer, 2016.

[29] Peng Zhang, Jiuling Wang, Ali Farhadi, Martial Hebert, and Devi Parikh. Predicting failures of vision systems. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 3566–3573, 2014.

[30] Shanshan Zhang, Rodrigo Benenson, Mohamed Omran, Jan Hosang, and Bernt Schiele. How far are we from solving pedestrian detection? In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 1259–1267, 2016.

[31] Shanshan Zhang, Rodrigo Benenson, and Bernt Schiele. Citypersons: A diverse dataset for pedestrian detection. In *CVPR*, 2017.